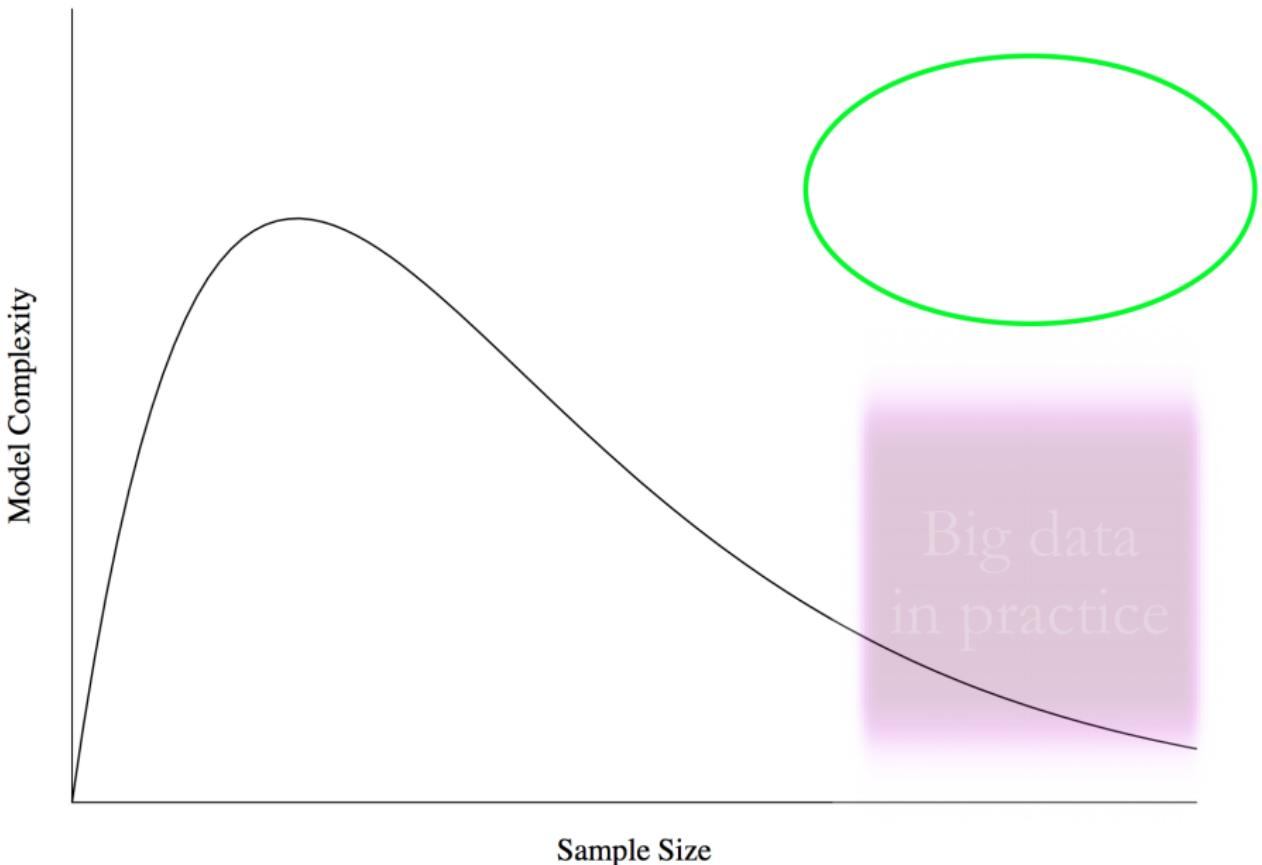


Big Data need Big Model



Stan: A platform for Bayesian inference

Andrew Gelman, Bob Carpenter, Matt Hoffman, Daniel Lee,
Ben Goodrich, Michael Betancourt, Marcus Brubaker,
Jiqiang Guo, Peter Li, Allen Riddell, . . .

Department of Statistics, Columbia University, New York
(and other places)

10 Nov 2014



Web Images Videos Maps News More ▾ Search tools

About 91,800,000 results (0.40 seconds)



[Eminem - Stan \(Short Version\) ft. Dido - YouTube](#)

www.youtube.com/watch?v=aSLZFdqwh7E ▾

Artists: Eminem, Dido

Album: No Angel





Eminem - Stan (Short Version) ft. Dido - YouTube

www.youtube.com/watch?v=aSLZFfdqwh7E ▾

Artists: Eminem, Dido

Album: No Angel

Released: 1999

[Feedback](#)

Stan (song) - Wikipedia, the free encyclopedia

[en.wikipedia.org/wiki/Stan_\(song\)](http://en.wikipedia.org/wiki/Stan_(song)) ▾

"Stan" is the third single from the The Marshall Mathers LP, recorded in 1999 by American rapper Eminem and featuring British singer Dido. It peaked at number ...

[Thank You - Rock City - Robert Browning - Murder ballad](#)

Stan: Project Home Page

mc-stan.org/ ▾

Stan modeling language and C++ library for Bayesian inference. NUTS adaptive HMC (MCMC) sampling, automatic differentiation, R, shell interfaces. Gelman.



Urban Dictionary: stan

www.urbandictionary.com/define.php?term=stan ▾

Based on the central character in the Eminem song of the same name, a "stan" is an overzealous maniacal fan for any celebrity or athlete.



Stan

Stan is a probabilistic programming language implementing full Bayesian statistical inference with

- MCMC sampling (NUTS, HMC)

and penalized maximum likelihood estimation with

- Optimization (BFGS)

Stan is coded in C++ and runs on all major platforms (Linux, Mac, Windows).

Stan is freedom-respecting, open-source software (new BSD core, GPLv3 interfaces).

Interfaces

Download and getting started instructions, organized by interface:

- [RStan v2.5.0](#) (R)
- [PyStan v2.5.0](#) (Python)
- [CmdStan v2.5.0](#) (shell, command-line terminal)
- [MatlabStan](#) (MATLAB)
- [Stan.jl](#) (Julia)

[Home](#)[RStan](#)[PyStan](#)[CmdStan](#)[MatlabStan](#)[Stan.jl](#)[Manual](#)[Examples](#)[Groups](#)[Issues](#)[Contribute](#)[Source](#)

Ordered probit

```
data {  
    int<lower=2> K;  
    int<lower=0> N;  
    int<lower=1> D;  
    int<lower=1,upper=K> y[N];  
    row_vector[D] x[N]; }  
parameters {  
    vector[D] beta;  
    ordered[K-1] c; }  
model {  
    vector[K] theta;  
    for (n in 1:N) {  
        real eta;  
        eta <- x[n] * beta;  
        theta[1] <- 1 - Phi(eta - c[1]);  
        for (k in 2:(K-1))  
            theta[k] <- Phi(eta - c[k-1]) - Phi(eta - c[k]);  
        theta[K] <- Phi(eta - c[K-1]);  
        y[n] ~ categorical(theta);  
    } }
```



Measurement error model

```
data {  
  ...  
  real x_meas[N];      // measurement of x  
  real<lower=0> tau;  // measurement noise  
}  
parameters {  
  real x[N];          // unknown true value  
  real mu_x;           // prior location  
  real sigma_x;        // prior scale  
  ...  
}  
model {  
  x ~ normal(mu_x, sigma_x); // prior  
  x_meas ~ normal(x, tau);   // measurement model  
  y ~ normal(alpha + beta * x, sigma);  
  ...  
}
```



Stan overview



Stan overview

- ▶ Fit open-ended Bayesian models



Stan overview

- ▶ Fit open-ended Bayesian models
- ▶ Specify log posterior density in C++



Stan overview

- ▶ Fit open-ended Bayesian models
- ▶ Specify log posterior density in C++
- ▶ Code a distribution once, then use it everywhere



Stan overview

- ▶ Fit open-ended Bayesian models
- ▶ Specify log posterior density in C++
- ▶ Code a distribution once, then use it everywhere
- ▶ Hamiltonian No-U-Turn sampler



Stan overview

- ▶ Fit open-ended Bayesian models
- ▶ Specify log posterior density in C++
- ▶ Code a distribution once, then use it everywhere
- ▶ Hamiltonian No-U-Turn sampler
- ▶ Autodiff



Stan overview

- ▶ Fit open-ended Bayesian models
- ▶ Specify log posterior density in C++
- ▶ Code a distribution once, then use it everywhere
- ▶ Hamiltonian No-U-Turn sampler
- ▶ Autodiff
- ▶ Runs from R, Python, Matlab, Julia; postprocessing



People



9/44

People

- ▶ Stan core (15)



People

- ▶ Stan core (15)
- ▶ Research collaborators (30)
- ▶ Developers (100)



People

- ▶ Stan core (15)
- ▶ Research collaborators (30)
- ▶ Developers (100)
- ▶ User community (1000)



People

- ▶ Stan core (15)
- ▶ Research collaborators (30)
- ▶ Developers (100)
- ▶ User community (1000)
- ▶ Users (10000)



Funding

- ▶ National Science Foundation
- ▶ Institute for Education Sciences
- ▶ Department of Energy
- ▶ Novartis
- ▶ YouGov



Roles of Stan



Roles of Stan

- ▶ Bayesian inference for unsophisticated users (alternative to Stata, Bugs, etc.)



Roles of Stan

- ▶ Bayesian inference for unsophisticated users (alternative to Stata, Bugs, etc.)
- ▶ Bayesian inference for sophisticated users (alternative to programming it yourself)



Roles of Stan

- ▶ Bayesian inference for unsophisticated users (alternative to Stata, Bugs, etc.)
- ▶ Bayesian inference for sophisticated users (alternative to programming it yourself)
- ▶ Fast and scalable gradient computation



Roles of Stan

- ▶ Bayesian inference for unsophisticated users (alternative to Stata, Bugs, etc.)
- ▶ Bayesian inference for sophisticated users (alternative to programming it yourself)
- ▶ Fast and scalable gradient computation
- ▶ Environment for developing new algorithms



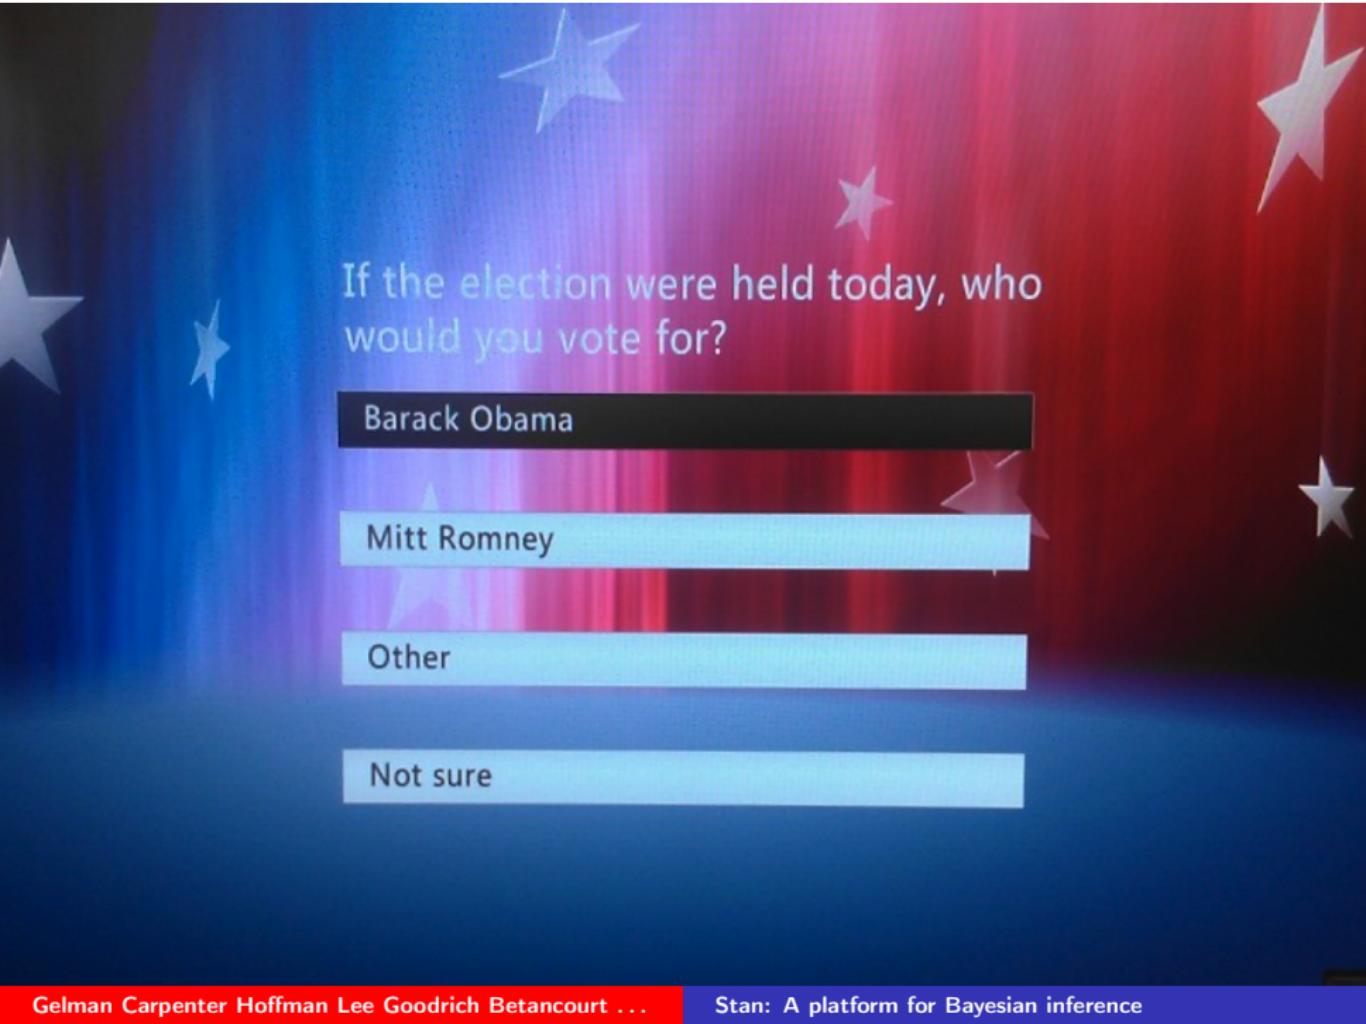
me social tv video games music apps



UFC
153

Press Conf





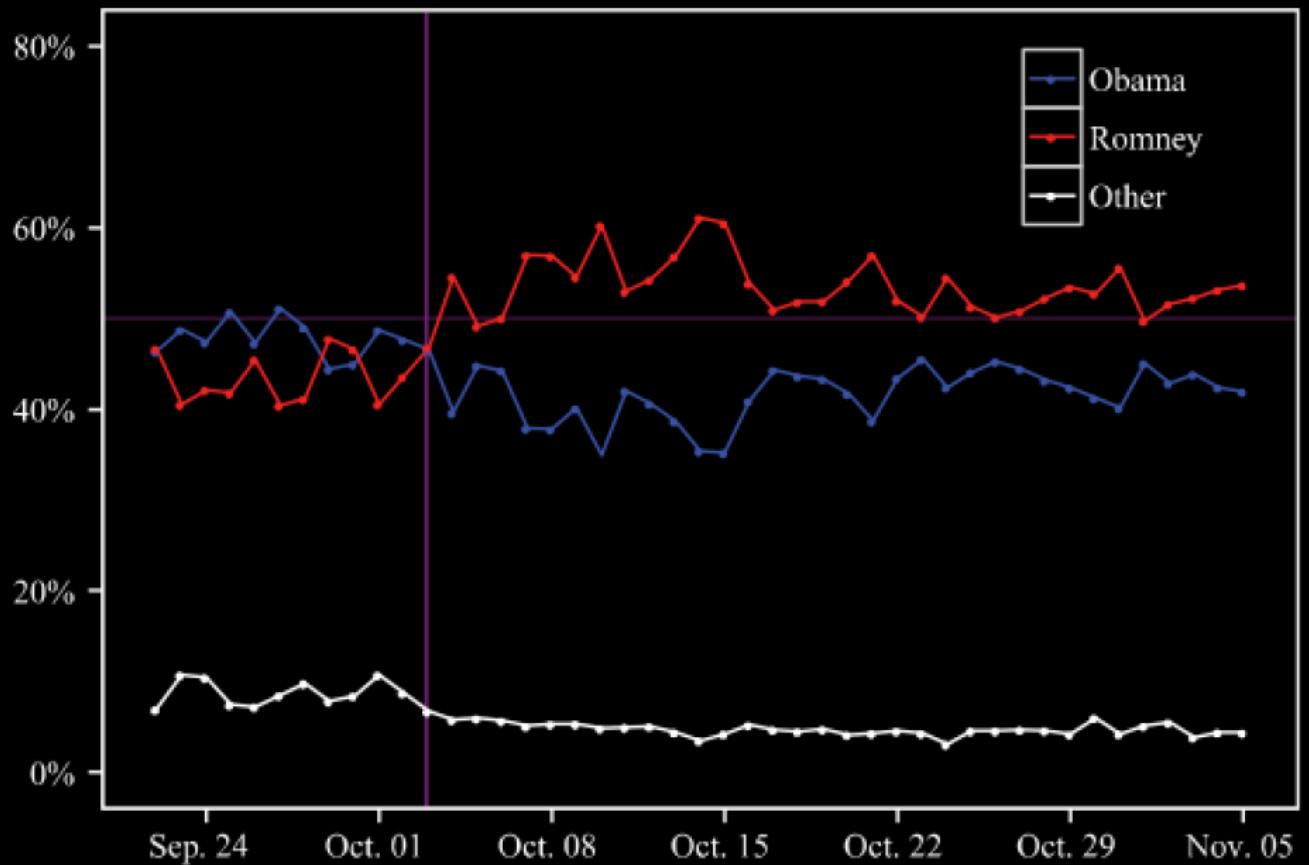
If the election were held today, who would you vote for?

Barack Obama

Mitt Romney

Other

Not sure



"This week, the New York Times and CBS News published a story using, in part, information from a non-probability, opt-in survey sparking concern among many in the polling community. In general, these methods have little grounding in theory and the results can vary widely based on the particular method used."

— Michael Link,
President, American Association for Public Opinion Research



Michael W. Link is Chief Methodologist for Research Methods at The Nielsen Company base of experience in survey research, having worked in academia (University of South Carolina, 1999), not-for-profit research (RTI International, 1999-2004), government (Centers for Disease Prevention, 2004-2007), and the private sector (Nielsen, 2007-present). He received his Science from the University of South Carolina. Michael's research centers around developing methodologies for confronting some of the most pressing issues facing survey research, techniques for improving survey participation and data quality (use of address-based sampling technologies), methodological issues involving use of multiple modes in data collection (mail, CATI, field, mobile, meters), and obtaining participation from hard-to-survey populations (isolated, racial and ethnic groups). His numerous research articles have appeared in *Public Opinion Quarterly* and other leading scientific journals.

An AAPOR member since 1993, Michael served as AAPOR Conference Chair in back-to-back years (2009 & 2010), a member of both the Cell Phone and Online task forces, an instructor for an AAPOR short course, a member of the Student Paper Competition Committee, a regular reviewer for *Public Opinion Quarterly*. He is a member of SAPOR, serving from 2009-2011 as President, Conference Chair, and Student Paper Competition Organizer and also a member of the Executive Committee.

In 2011 he, along with several research colleagues, received AAPOR's Warren J. Mitofsky Award for their work on address based sampling designs. His current research focuses on new technologies, such as mobile and social platforms, as vehicles for measuring and understanding attitudes and behaviors. He will be teaching a short course on "The Role of New Technologies in Augmenting, or Replacing Traditional Surveys" at the 2012 AAPOR conference.

Nielsen feels the heat of competition as it flubs its ratings of news broadcasts, putting ABC ahead of NBC

In spite of the goof, its global president took time to slam rival Rentrak, which collects different kind of data from viewers

NEW YORK DAILY NEWS / Sunday, October 19, 2014, 2:00 AM

BY DON KAPLAN



MEDIA

TV Ratings by Nielsen Had Errors for Months

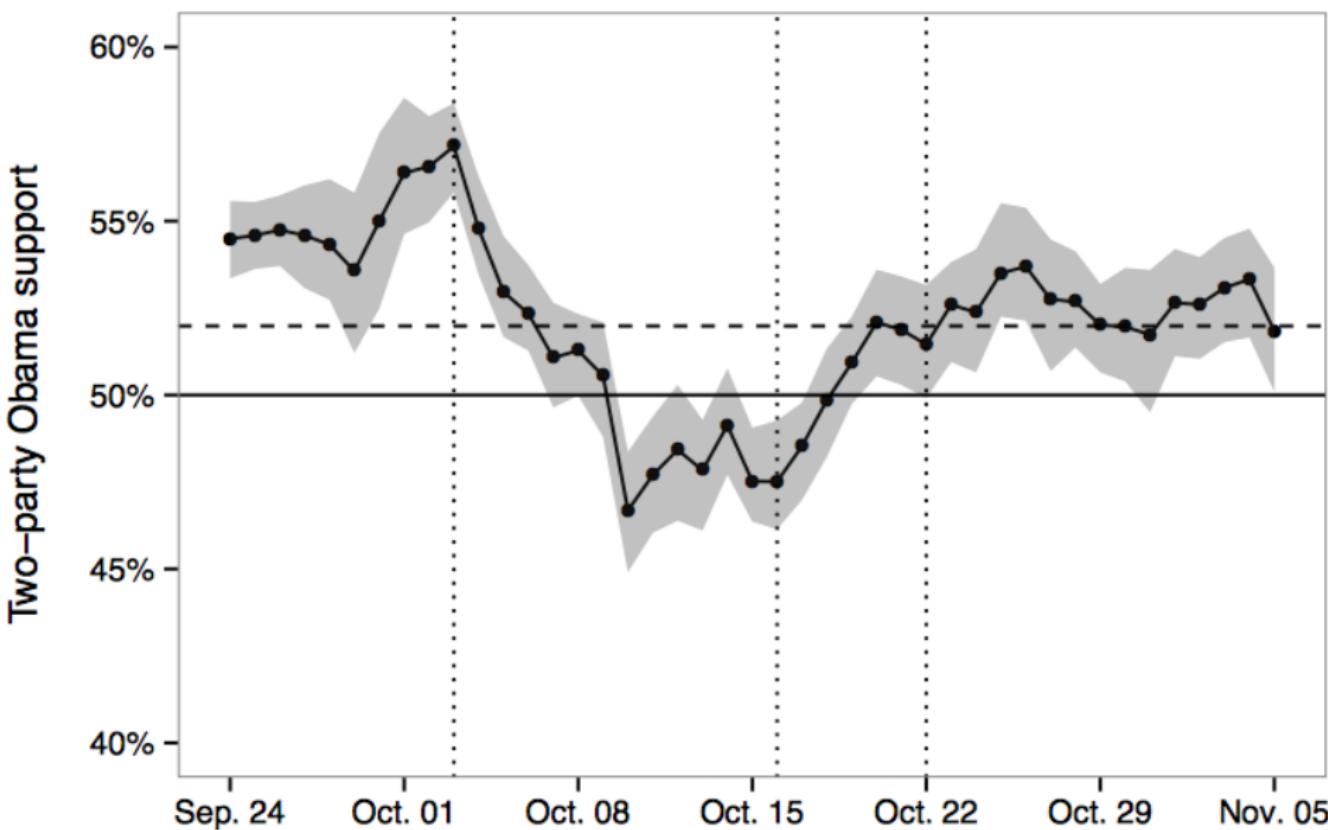
By BILL CARTER and EMILY STEEL OCT. 10, 2014

Email

Share

Nielsen, the television research firm, acknowledged on Friday that it had been reporting inaccurate ratings for the broadcast networks for the last seven months, a mistake that raises questions about the company's increasingly criticized system for measuring TV audiences.

Xbox estimates, adjusting for demographics





Gelman Carpenter Hoffman Lee Goodrich Betancourt ...



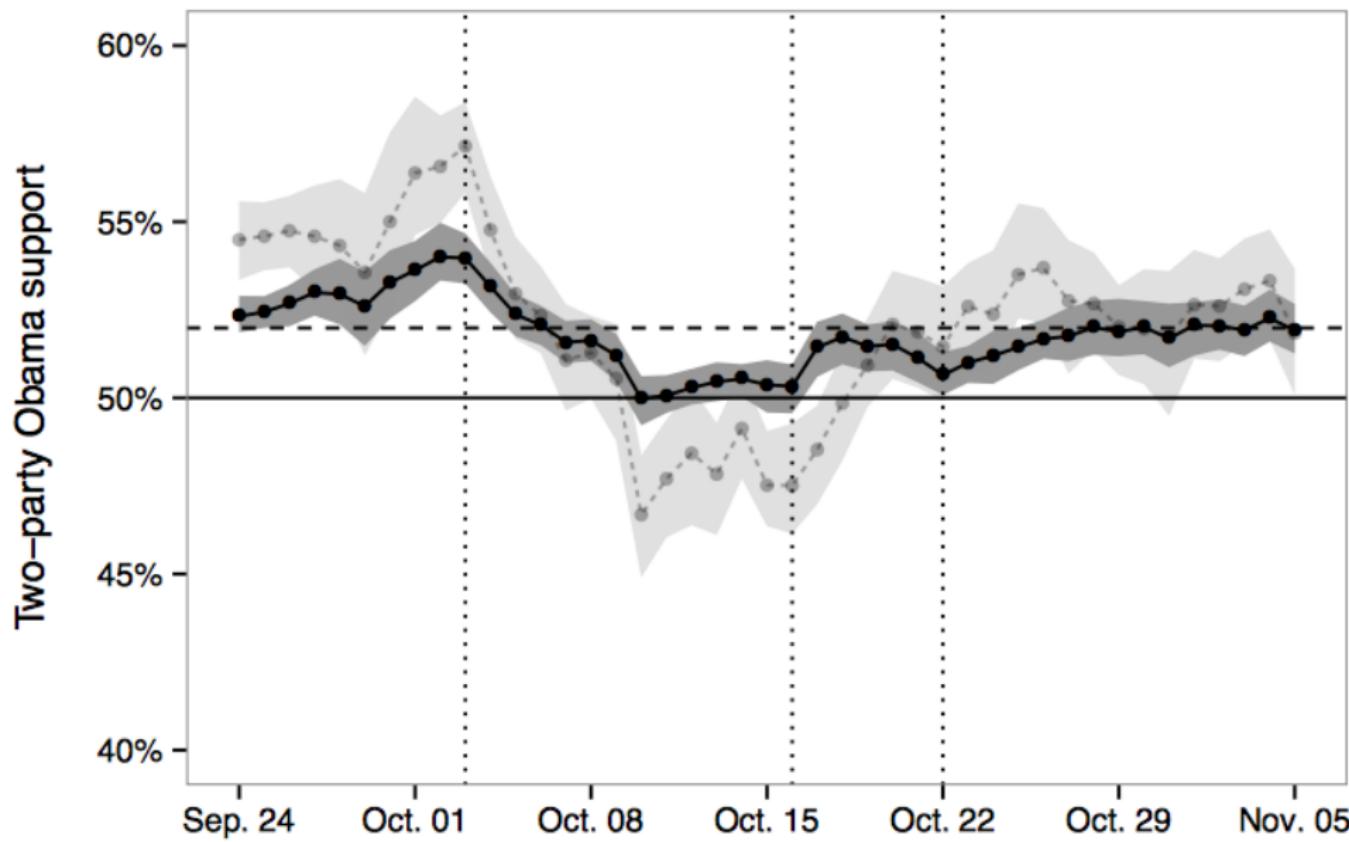
Stan: A platform for Bayesian inference



- ▶ Karl Rove, *Wall Street Journal*, 7 Oct: “Mr. Romney’s bounce is significant.”
- ▶ Nate Silver, *New York Times*, 6 Oct: “Mr. Romney has not only improved his own standing but also taken voters away from Mr. Obama’s column.”

STAN

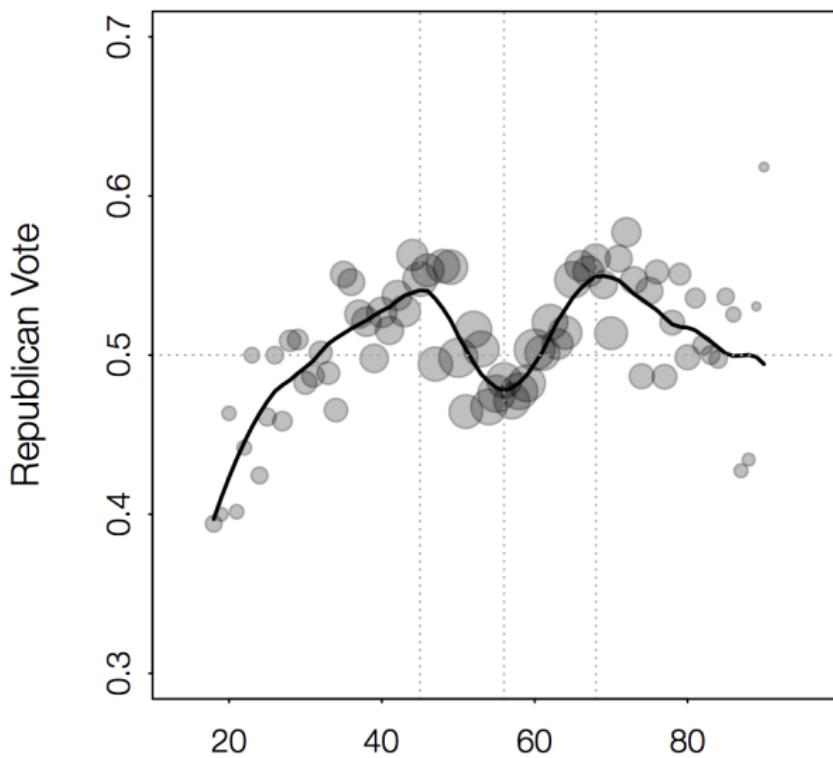
Xbox estimates, adjusting for demographics and partisanship



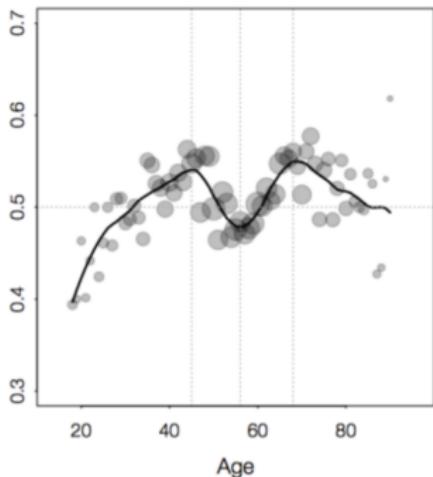
Jimmy Carter Republicans and George W. Bush Democrats



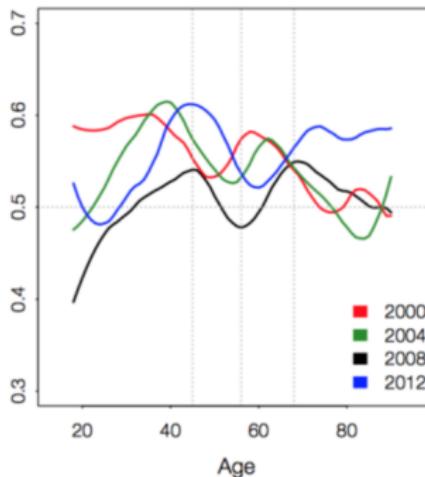
Non-Monotonic Age Curve in 2008



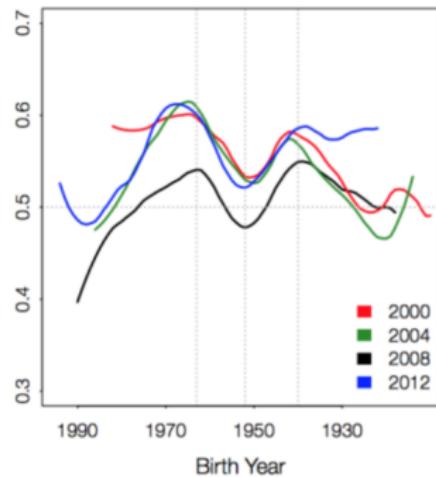
Non-Monotonic Age Curve in 2008



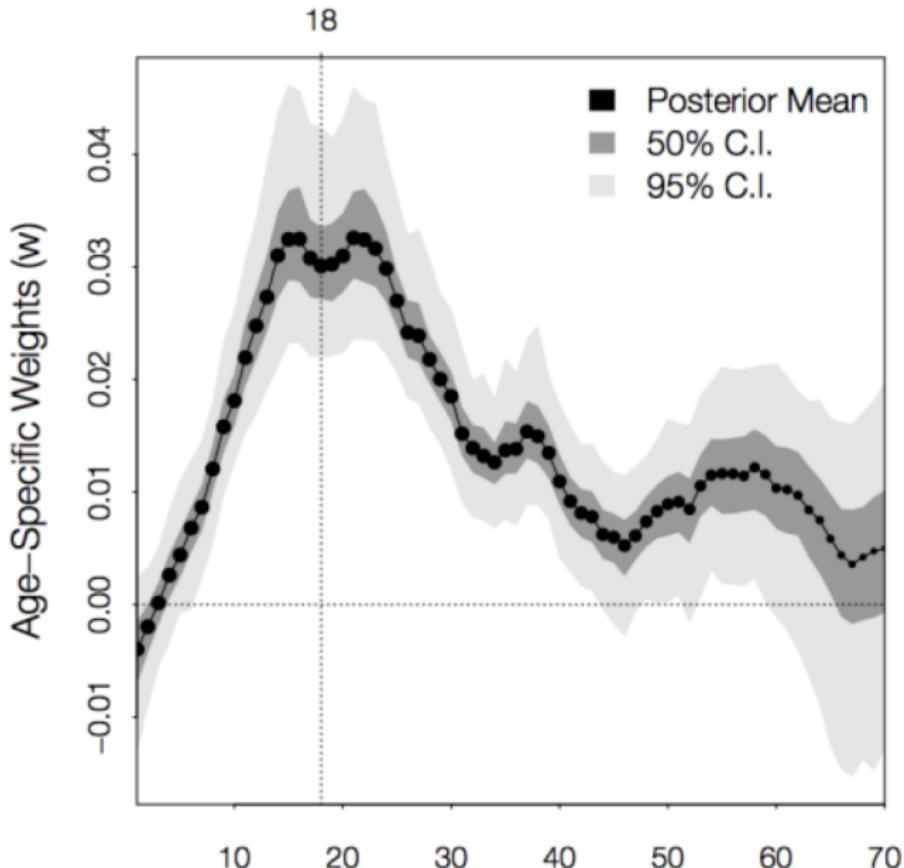
Non-Monotonicity in Other Elections



Lining up by Birth Year



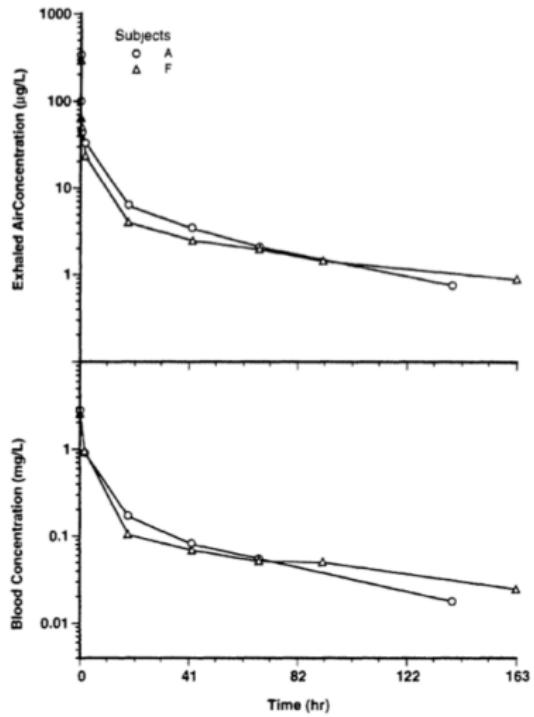
The Formative Years



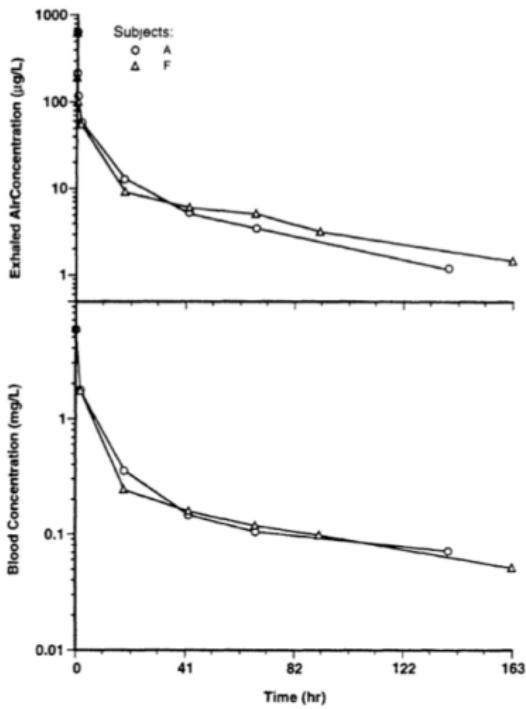
STAN

Toxicology

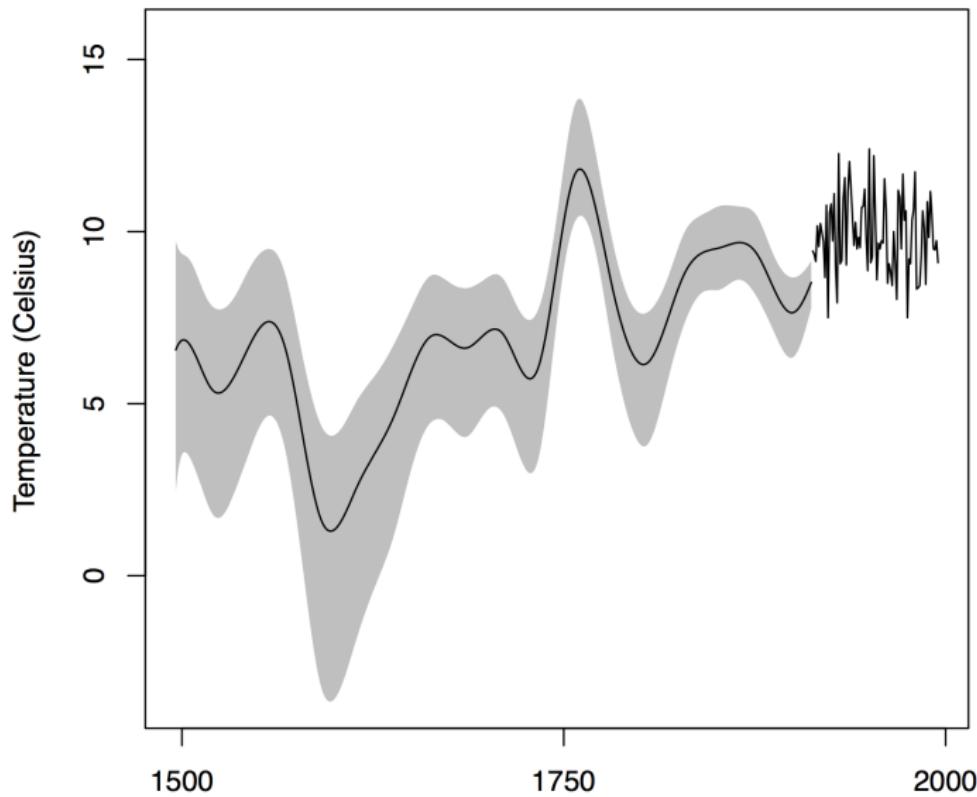
Exposure of 72 ppm



Exposure of 144 ppm



Earth science



Lots of other applications

Astronomy, ecology, linguistics, epidemiology, soil science, ...



Steps of Bayesian data analysis



Steps of Bayesian data analysis

- ▶ Model building



Steps of Bayesian data analysis

- ▶ Model building
- ▶ Inference



Steps of Bayesian data analysis

- ▶ Model building
- ▶ Inference
- ▶ Model checking



Steps of Bayesian data analysis

- ▶ Model building
- ▶ Inference
- ▶ Model checking
- ▶ Model understanding and improvement



Background on Bayesian computation



Background on Bayesian computation

- ▶ Point estimates and standard errors



Background on Bayesian computation

- ▶ Point estimates and standard errors
- ▶ Hierarchical models



Background on Bayesian computation

- ▶ Point estimates and standard errors
- ▶ Hierarchical models
- ▶ Posterior simulation



Background on Bayesian computation

- ▶ Point estimates and standard errors
- ▶ Hierarchical models
- ▶ Posterior simulation
- ▶ Markov chain Monte Carlo (Gibbs sampler and Metropolis algorithm)



Background on Bayesian computation

- ▶ Point estimates and standard errors
- ▶ Hierarchical models
- ▶ Posterior simulation
- ▶ Markov chain Monte Carlo (Gibbs sampler and Metropolis algorithm)
- ▶ Hamiltonian Monte Carlo



Solving problems



Solving problems

- ▶ *Problem:* Gibbs too slow, Metropolis too problem-specific



Solving problems

- ▶ *Problem:* Gibbs too slow, Metropolis too problem-specific
- ▶ *Solution:* Hamiltonian Monte Carlo



Solving problems

- ▶ *Problem:* Gibbs too slow, Metropolis too problem-specific
- ▶ *Solution:* Hamiltonian Monte Carlo
- ▶ *Problem:* Interpreters too slow, won't scale



Solving problems

- ▶ *Problem:* Gibbs too slow, Metropolis too problem-specific
- ▶ *Solution:* Hamiltonian Monte Carlo
- ▶ *Problem:* Interpreters too slow, won't scale
- ▶ *Solution:* Compilation



Solving problems

- ▶ *Problem:* Gibbs too slow, Metropolis too problem-specific
- ▶ *Solution:* Hamiltonian Monte Carlo
- ▶ *Problem:* Interpreters too slow, won't scale
- ▶ *Solution:* Compilation
- ▶ *Problem:* Need gradients of log posterior for HMC



Solving problems

- ▶ *Problem:* Gibbs too slow, Metropolis too problem-specific
- ▶ *Solution:* Hamiltonian Monte Carlo
- ▶ *Problem:* Interpreters too slow, won't scale
- ▶ *Solution:* Compilation
- ▶ *Problem:* Need gradients of log posterior for HMC
- ▶ *Solution:* Reverse-mode algorithmic differentiation



Solving problems

- ▶ *Problem:* Gibbs too slow, Metropolis too problem-specific
- ▶ *Solution:* Hamiltonian Monte Carlo

- ▶ *Problem:* Interpreters too slow, won't scale
- ▶ *Solution:* Compilation

- ▶ *Problem:* Need gradients of log posterior for HMC
- ▶ *Solution:* Reverse-mode algorithmic differentiation

- ▶ *Problem:* Existing algo-diff slow, limited, unextensible



Solving problems

- ▶ *Problem:* Gibbs too slow, Metropolis too problem-specific
- ▶ *Solution:* Hamiltonian Monte Carlo

- ▶ *Problem:* Interpreters too slow, won't scale
- ▶ *Solution:* Compilation

- ▶ *Problem:* Need gradients of log posterior for HMC
- ▶ *Solution:* Reverse-mode algorithmic differentiation

- ▶ *Problem:* Existing algo-diff slow, limited, unextensible
- ▶ *Solution:* Our own algo-diff



Solving problems

- ▶ *Problem:* Gibbs too slow, Metropolis too problem-specific
- ▶ *Solution:* Hamiltonian Monte Carlo

- ▶ *Problem:* Interpreters too slow, won't scale
- ▶ *Solution:* Compilation

- ▶ *Problem:* Need gradients of log posterior for HMC
- ▶ *Solution:* Reverse-mode algorithmic differentiation

- ▶ *Problem:* Existing algo-diff slow, limited, unextensible
- ▶ *Solution:* Our own algo-diff

- ▶ *Problem:* Algo-diff requires fully templated functions



Solving problems

- ▶ *Problem:* Gibbs too slow, Metropolis too problem-specific
- ▶ *Solution:* Hamiltonian Monte Carlo

- ▶ *Problem:* Interpreters too slow, won't scale
- ▶ *Solution:* Compilation

- ▶ *Problem:* Need gradients of log posterior for HMC
- ▶ *Solution:* Reverse-mode algorithmic differentiation

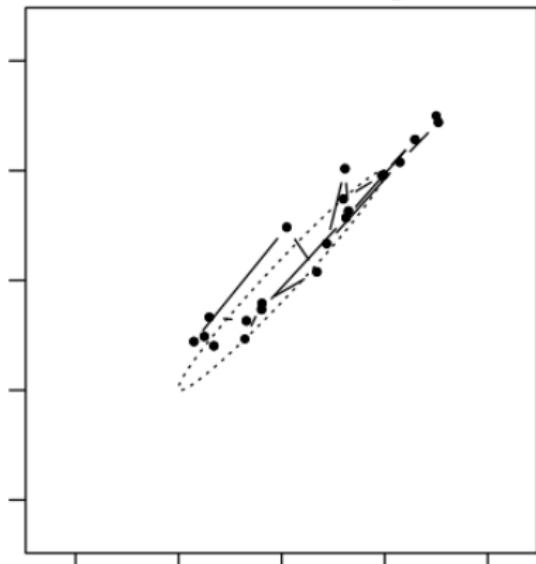
- ▶ *Problem:* Existing algo-diff slow, limited, unextensible
- ▶ *Solution:* Our own algo-diff

- ▶ *Problem:* Algo-diff requires fully templated functions
- ▶ *Solution:* Our own density library, Eigen linear algebra

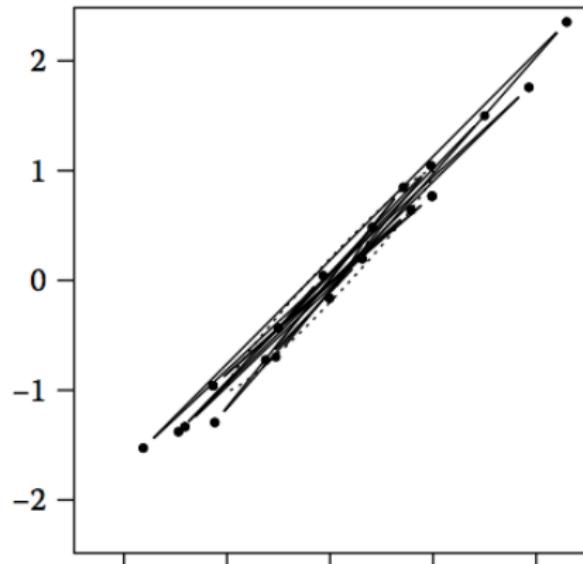


Radford Neal (2011) on Hamiltonian Monte Carlo

Random-walk Metropolis



Hamiltonian Monte Carlo



"One practical impediment to the use of Hamiltonian Monte Carlo is the need to select suitable values for the leapfrog stepsize, ϵ , and the number of leapfrog steps L . . . Tuning HMC will usually require preliminary runs with trial values for ϵ and L . . . Unfortunately, preliminary runs can be misleading . . . "

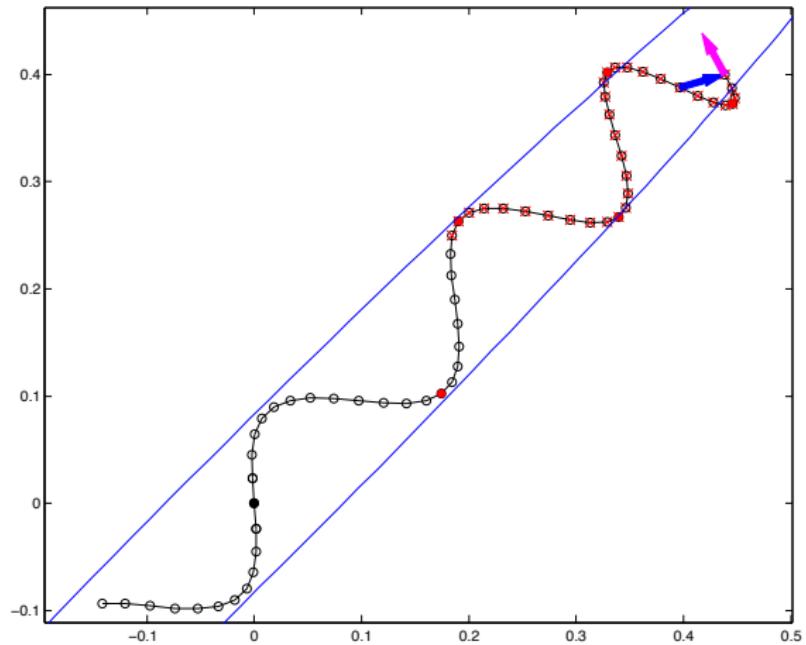


The No U-Turn Sampler

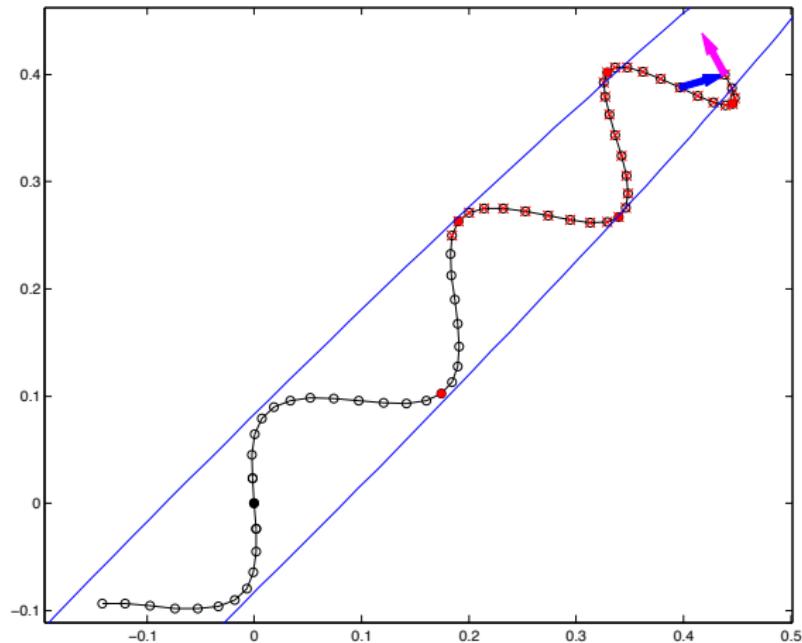
- ▶ Created by Matt Hoffman
- ▶ Run the HMC steps until they start to turn around
(bend with an angle $> 180^\circ$)
- ▶ Computationally efficient
- ▶ Requires no tuning of #steps
- ▶ Complications to preserve detailed balance



NUTS Example Trajectory



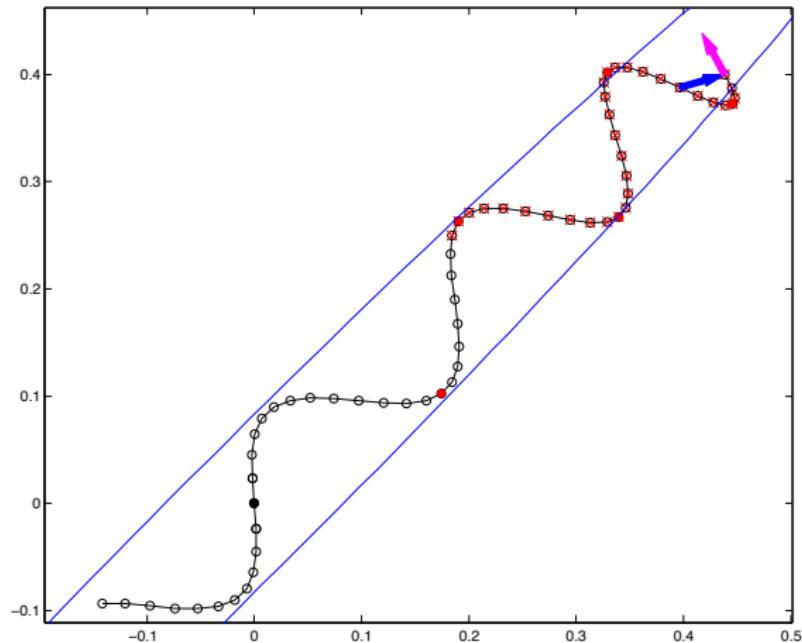
NUTS Example Trajectory



- ▶ Blue ellipse is contour of target distribution



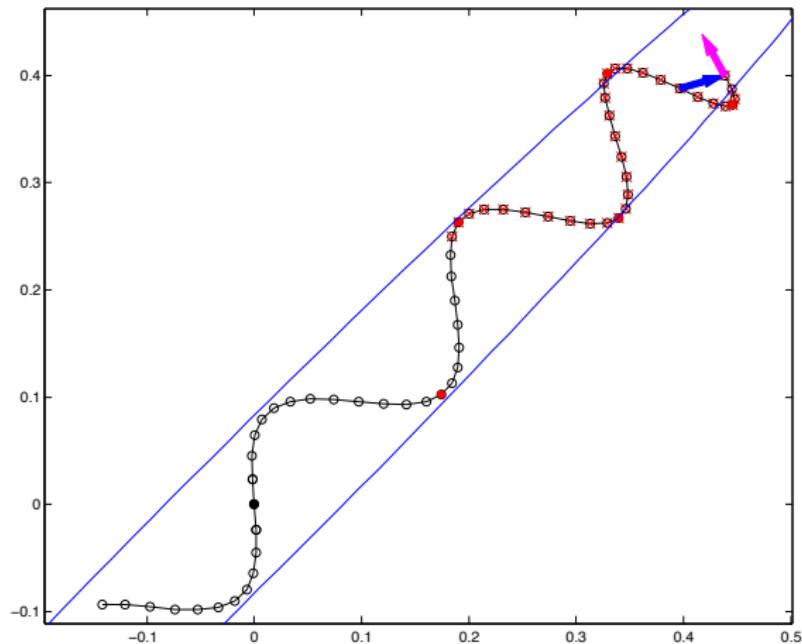
NUTS Example Trajectory



- ▶ Blue ellipse is contour of target distribution
- ▶ Initial position at black solid circle



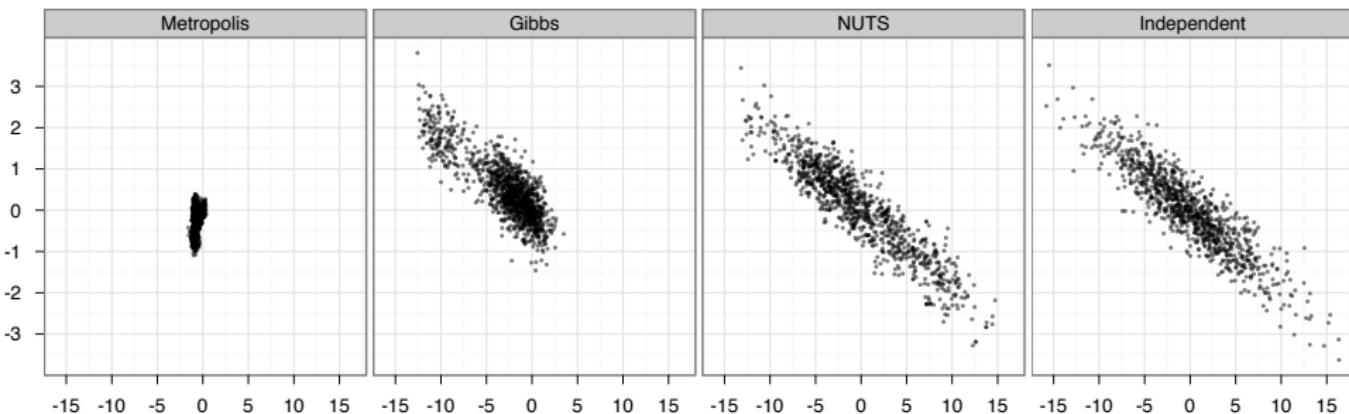
NUTS Example Trajectory



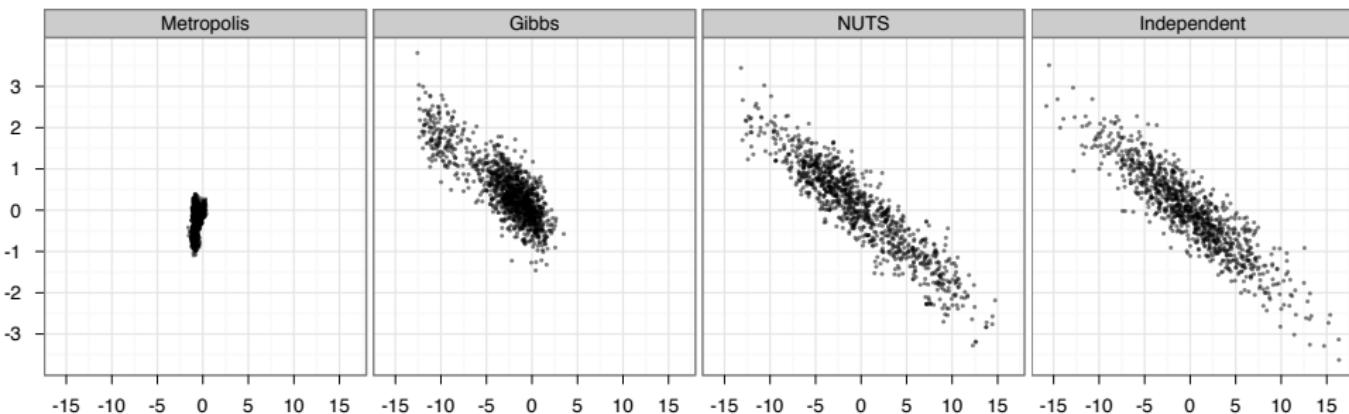
- ▶ Blue ellipse is contour of target distribution
- ▶ Initial position at black solid circle
- ▶ Arrows indicate a U-turn in momentum



NUTS vs. Gibbs and Metropolis



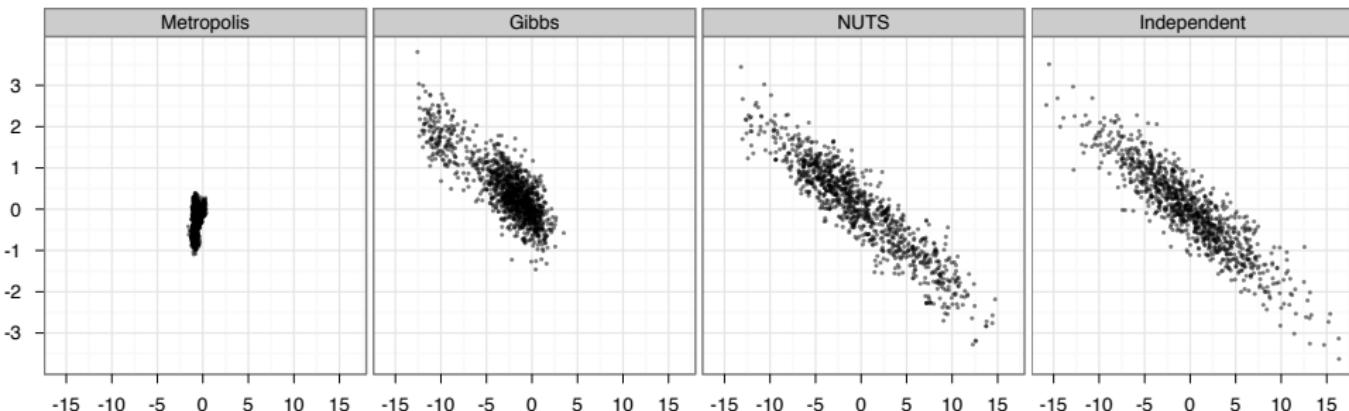
NUTS vs. Gibbs and Metropolis



- ▶ Two dimensions of highly correlated 250-dim distribution



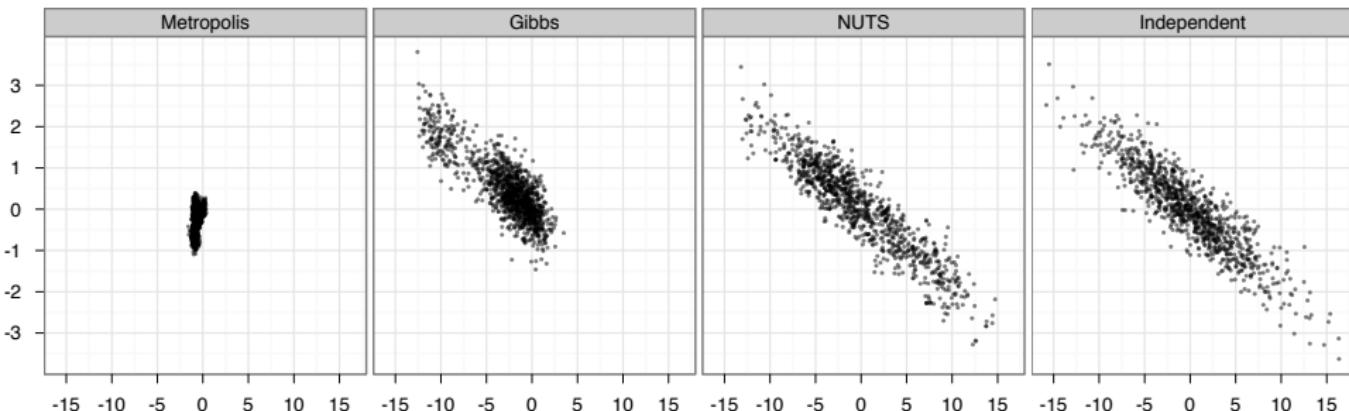
NUTS vs. Gibbs and Metropolis



- ▶ Two dimensions of highly correlated 250-dim distribution
- ▶ 1M samples from Metropolis, 1M from Gibbs (thin to 1K)



NUTS vs. Gibbs and Metropolis

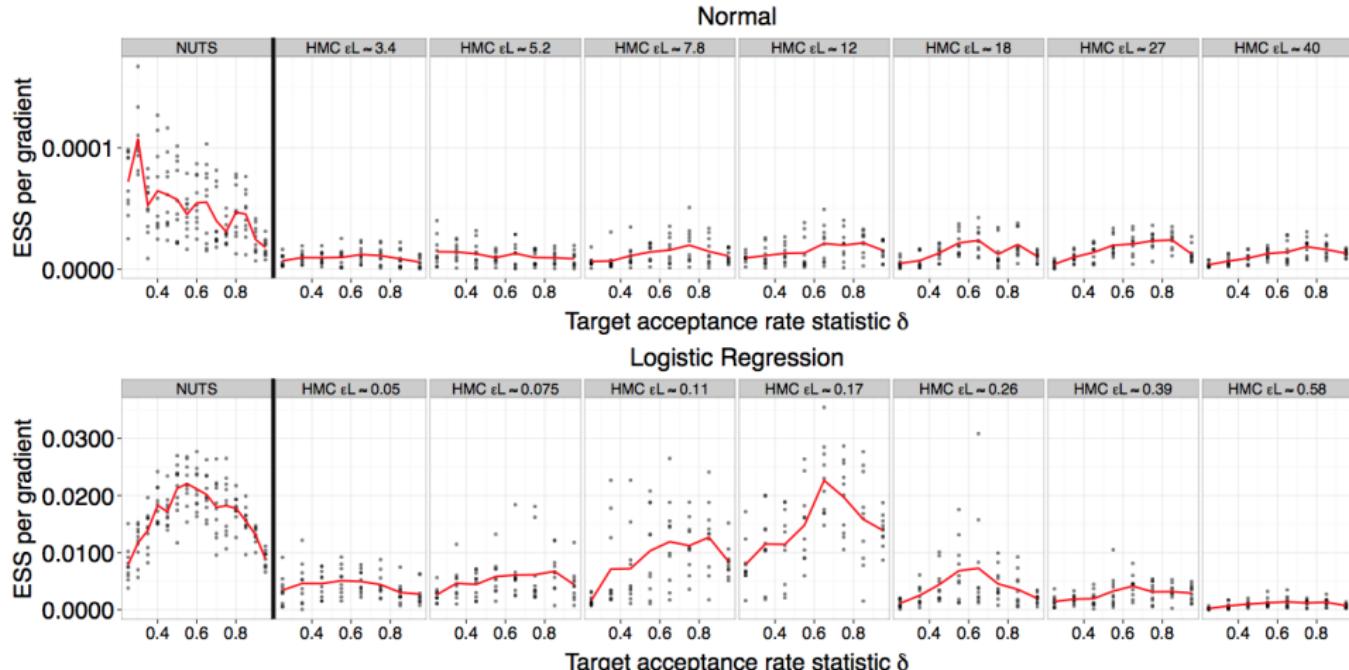


- ▶ Two dimensions of highly correlated 250-dim distribution
- ▶ 1M samples from Metropolis, 1M from Gibbs (thin to 1K)
- ▶ 1K samples from NUTS, 1K independent draws



NUTS vs. Basic HMC

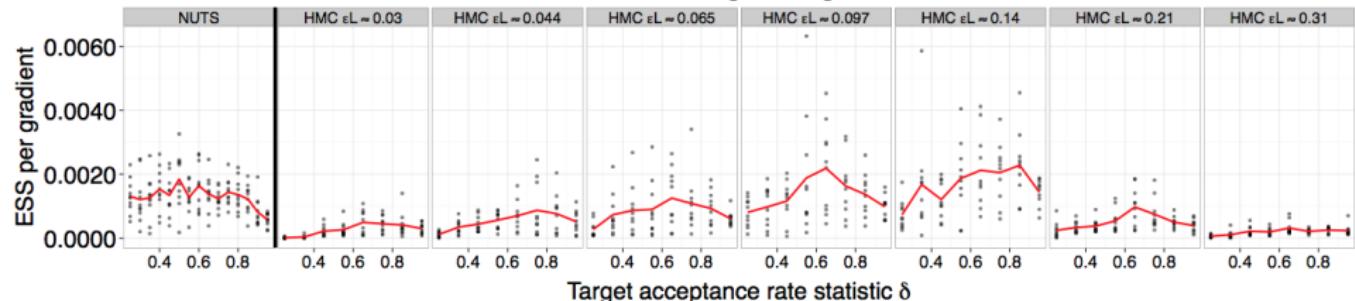
- ▶ 250-D normal and logistic regression models
- ▶ Vertical axis shows effective #sims (big is good)
- ▶ (Left) NUTS; (Right) HMC with increasing $t = \epsilon L$



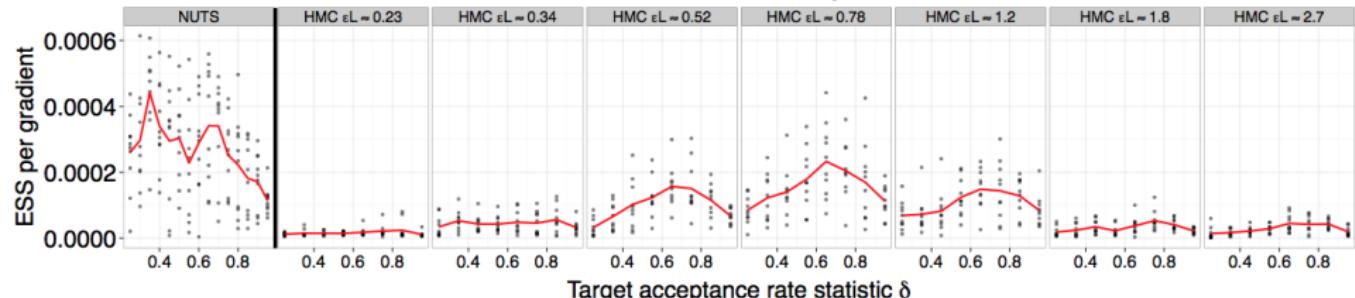
NUTS vs. Basic HMC II

- ▶ Hierarchical logistic regression and stochastic volatility
- ▶ Simulation time is step size ϵ times #steps L
- ▶ NUTS can beat optimally tuned HMC

Hierarchical Logistic Regression



Stochastic Volatility



Solving more problems in Stan



Solving more problems in Stan

- ▶ *Problem:* Need ease of use of BUGS



Solving more problems in Stan

- ▶ *Problem:* Need ease of use of BUGS
- ▶ *Solution:* Compile directed graphical model language



Solving more problems in Stan

- ▶ *Problem:* Need ease of use of BUGS
- ▶ *Solution:* Compile directed graphical model language
- ▶ *Problem:* Need to tune parameters for HMC



Solving more problems in Stan

- ▶ *Problem:* Need ease of use of BUGS
- ▶ *Solution:* Compile directed graphical model language
- ▶ *Problem:* Need to tune parameters for HMC
- ▶ *Solution:* Auto tuning, adaptation



Solving more problems in Stan

- ▶ *Problem:* Need ease of use of BUGS
- ▶ *Solution:* Compile directed graphical model language
- ▶ *Problem:* Need to tune parameters for HMC
- ▶ *Solution:* Auto tuning, adaptation
- ▶ *Problem:* Efficient up-to-proportion density calcs



Solving more problems in Stan

- ▶ *Problem:* Need ease of use of BUGS
- ▶ *Solution:* Compile directed graphical model language

- ▶ *Problem:* Need to tune parameters for HMC
- ▶ *Solution:* Auto tuning, adaptation

- ▶ *Problem:* Efficient up-to-proportion density calcs
- ▶ *Solution:* Density template metaprogramming



Solving more problems in Stan

- ▶ *Problem:* Need ease of use of BUGS
- ▶ *Solution:* Compile directed graphical model language

- ▶ *Problem:* Need to tune parameters for HMC
- ▶ *Solution:* Auto tuning, adaptation

- ▶ *Problem:* Efficient up-to-proportion density calcs
- ▶ *Solution:* Density template metaprogramming

- ▶ *Problem:* Limited error checking, recovery



Solving more problems in Stan

- ▶ *Problem:* Need ease of use of BUGS
- ▶ *Solution:* Compile directed graphical model language

- ▶ *Problem:* Need to tune parameters for HMC
- ▶ *Solution:* Auto tuning, adaptation

- ▶ *Problem:* Efficient up-to-proportion density calcs
- ▶ *Solution:* Density template metaprogramming

- ▶ *Problem:* Limited error checking, recovery
- ▶ *Solution:* Static model typing, informative exceptions



Solving more problems in Stan

- ▶ *Problem:* Need ease of use of BUGS
- ▶ *Solution:* Compile directed graphical model language

- ▶ *Problem:* Need to tune parameters for HMC
- ▶ *Solution:* Auto tuning, adaptation

- ▶ *Problem:* Efficient up-to-proportion density calcs
- ▶ *Solution:* Density template metaprogramming

- ▶ *Problem:* Limited error checking, recovery
- ▶ *Solution:* Static model typing, informative exceptions

- ▶ *Problem:* Poor boundary behavior



Solving more problems in Stan

- ▶ *Problem:* Need ease of use of BUGS
- ▶ *Solution:* Compile directed graphical model language

- ▶ *Problem:* Need to tune parameters for HMC
- ▶ *Solution:* Auto tuning, adaptation

- ▶ *Problem:* Efficient up-to-proportion density calcs
- ▶ *Solution:* Density template metaprogramming

- ▶ *Problem:* Limited error checking, recovery
- ▶ *Solution:* Static model typing, informative exceptions

- ▶ *Problem:* Poor boundary behavior
- ▶ *Solution:* Calculate limits (e.g. $\lim_{x \rightarrow 0} x \log x$)



Solving more problems in Stan

- ▶ *Problem:* Need ease of use of BUGS
- ▶ *Solution:* Compile directed graphical model language

- ▶ *Problem:* Need to tune parameters for HMC
- ▶ *Solution:* Auto tuning, adaptation

- ▶ *Problem:* Efficient up-to-proportion density calcs
- ▶ *Solution:* Density template metaprogramming

- ▶ *Problem:* Limited error checking, recovery
- ▶ *Solution:* Static model typing, informative exceptions

- ▶ *Problem:* Poor boundary behavior
- ▶ *Solution:* Calculate limits (e.g. $\lim_{x \rightarrow 0} x \log x$)



Solving more problems in Stan

- ▶ *Problem:* Need ease of use of BUGS
- ▶ *Solution:* Compile directed graphical model language

- ▶ *Problem:* Need to tune parameters for HMC
- ▶ *Solution:* Auto tuning, adaptation

- ▶ *Problem:* Efficient up-to-proportion density calcs
- ▶ *Solution:* Density template metaprogramming

- ▶ *Problem:* Limited error checking, recovery
- ▶ *Solution:* Static model typing, informative exceptions

- ▶ *Problem:* Poor boundary behavior
- ▶ *Solution:* Calculate limits (e.g. $\lim_{x \rightarrow 0} x \log x$)

- ▶ *Problem:* Restrictive licensing (e.g., closed, GPL, etc.)



Solving more problems in Stan

- ▶ *Problem:* Need ease of use of BUGS
- ▶ *Solution:* Compile directed graphical model language

- ▶ *Problem:* Need to tune parameters for HMC
- ▶ *Solution:* Auto tuning, adaptation

- ▶ *Problem:* Efficient up-to-proportion density calcs
- ▶ *Solution:* Density template metaprogramming

- ▶ *Problem:* Limited error checking, recovery
- ▶ *Solution:* Static model typing, informative exceptions

- ▶ *Problem:* Poor boundary behavior
- ▶ *Solution:* Calculate limits (e.g. $\lim_{x \rightarrow 0} x \log x$)

- ▶ *Problem:* Restrictive licensing (e.g., closed, GPL, etc.)
- ▶ *Solution:* Open-source, BSD license



New stuff: Differential equation models

Simple harmonic oscillator:

$$\begin{aligned}\frac{dz_1}{dt} &= -z_2 \\ \frac{dz_2}{dt} &= -z_1 - \theta z_2\end{aligned}$$

with observations $(y_1, y_2)_t, t = 1, \dots, T$:

$$y_{1t} \sim N(z_1(t), \sigma_1^2)$$

$$y_{2t} \sim N(z_2(t), \sigma_2^2)$$

Given data $(y_1, y_2)_t, t = 1, \dots, T$,
estimate initial state $(y_1, y_2)_{t=0}$ and parameter θ



Stan program: 1

```
functions {
    real[] sho(real t, real[] y, real[] theta, real[] x_r, int[] x_i) {
        real dydt[2];
        dydt[1] <- y[2];
        dydt[2] <- -y[1] - theta[1] * y[2];
        return dydt;
    }
}
data {
    int<lower=1> T;
    real y[T,2];
    real t0;
    real ts[T];
}
transformed data {
    real x_r[0];
    int x_i[0];
}
```



Stan program: 2

```
parameters {
    real y0[2];
    vector<lower=0>[2] sigma;
    real theta[1];
}
model {
    real z[T,2];
    sigma ~ cauchy(0,2.5);
    theta ~ normal(0,1);
    y0 ~ normal(0,1);
    z <- integrate_ode(sho, y0, t0, ts, theta, x_r, x_i);
    for (t in 1:T)
        y[t] ~ normal(z[t], sigma);
}
```



Stan output

Run RStan with data simulated from
 $\theta = 0.15$, $y_0 = (1, 0)$, and $\sigma = 0.1$:

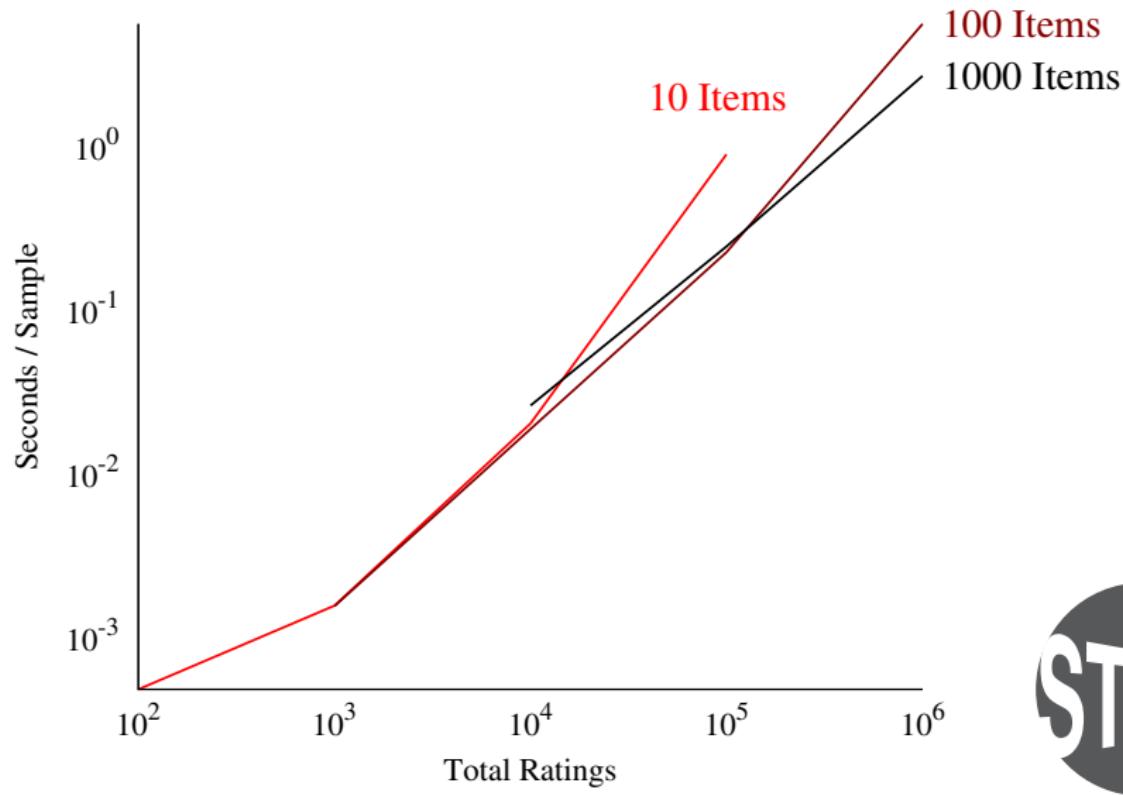
Inference for Stan model: sho.

4 chains, each with iter=2000; warmup=1000; thin=1;
post-warmup draws per chain=1000, total post-warmup draws=4000.

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
y0[1]	1.05	0.00	0.09	0.87	0.98	1.05	1.10	1.23	1172	1
y0[2]	-0.06	0.00	0.06	-0.18	-0.10	-0.06	-0.02	0.06	1524	1
sigma[1]	0.14	0.00	0.04	0.08	0.11	0.13	0.16	0.25	1354	1
sigma[2]	0.11	0.00	0.03	0.06	0.08	0.10	0.12	0.18	1697	1
theta[1]	0.15	0.00	0.04	0.08	0.13	0.15	0.17	0.22	1112	1
lp__	28.97	0.06	1.80	24.55	27.95	29.37	30.29	31.35	992	1



Big Data, Big Model, Scalable Computing



Thinking about scalability



Thinking about scalability

- ▶ Hierarchical item response model:

# items	# raters	# groups	# data	Stan		JAGS	
				time	memory	time	memory
20	2,000	100	40,000	:02m	16MB	:03m	220MB
40	8,000	200	320,000	:16m	92MB	:40m	1400MB
80	32,000	400	2,560,000	4h:10m	580MB	:??m	?MB



Thinking about scalability

- ▶ Hierarchical item response model:

# items	# raters	# groups	# data	Stan		JAGS	
				time	memory	time	memory
20	2,000	100	40,000	:02m	16MB	:03m	220MB
40	8,000	200	320,000	:16m	92MB	:40m	1400MB
80	32,000	400	2,560,000	4h:10m	580MB	:??m	?MB

- ▶ Also, Stan generated 4x effective sample size per iteration



Future work



Future work

- ▶ Programming



Future work

- ▶ Programming
 - ▶ Faster gradients and higher-order derivatives



Future work

- ▶ Programming
 - ▶ Faster gradients and higher-order derivatives
 - ▶ Functions



Future work

- ▶ Programming
 - ▶ Faster gradients and higher-order derivatives
 - ▶ Functions
- ▶ Statistical algorithms



Future work

- ▶ Programming
 - ▶ Faster gradients and higher-order derivatives
 - ▶ Functions
- ▶ Statistical algorithms
 - ▶ Riemannian Hamiltonian Monte Carlo



Future work

- ▶ Programming
 - ▶ Faster gradients and higher-order derivatives
 - ▶ Functions
- ▶ Statistical algorithms
 - ▶ Riemannian Hamiltonian Monte Carlo
 - ▶ (Penalized) mle



Future work

- ▶ Programming
 - ▶ Faster gradients and higher-order derivatives
 - ▶ Functions
- ▶ Statistical algorithms
 - ▶ Riemannian Hamiltonian Monte Carlo
 - ▶ (Penalized) mle
 - ▶ (Penalized) marginal mle



Future work

- ▶ Programming
 - ▶ Faster gradients and higher-order derivatives
 - ▶ Functions
- ▶ Statistical algorithms
 - ▶ Riemannian Hamiltonian Monte Carlo
 - ▶ (Penalized) mle
 - ▶ (Penalized) marginal mle
 - ▶ Black-box variational Bayes



Future work

- ▶ Programming
 - ▶ Faster gradients and higher-order derivatives
 - ▶ Functions
- ▶ Statistical algorithms
 - ▶ Riemannian Hamiltonian Monte Carlo
 - ▶ (Penalized) mle
 - ▶ (Penalized) marginal mle
 - ▶ Black-box variational Bayes
 - ▶ Data partitioning and expectation propagation

